



LAB 7:

# ADVANCED CNN ARCHITECTURES

University of Washington, Seattle

Spring 2025



# OUTLINE

## Part 1: Limitations of conventional CNNs

- Very deep networks
- Degradation problem

## Part 2: Additional CNN components

- Residual block
- Bottleneck block
- Inception modules

## Part 3: ResNet example

- MNIST classification

## Part 4: Additional resource: UW Hayak

- UW Hyak
- Access through RCC

## Part 5: Lab Assignment

- Cifar-10 classification with ResNet



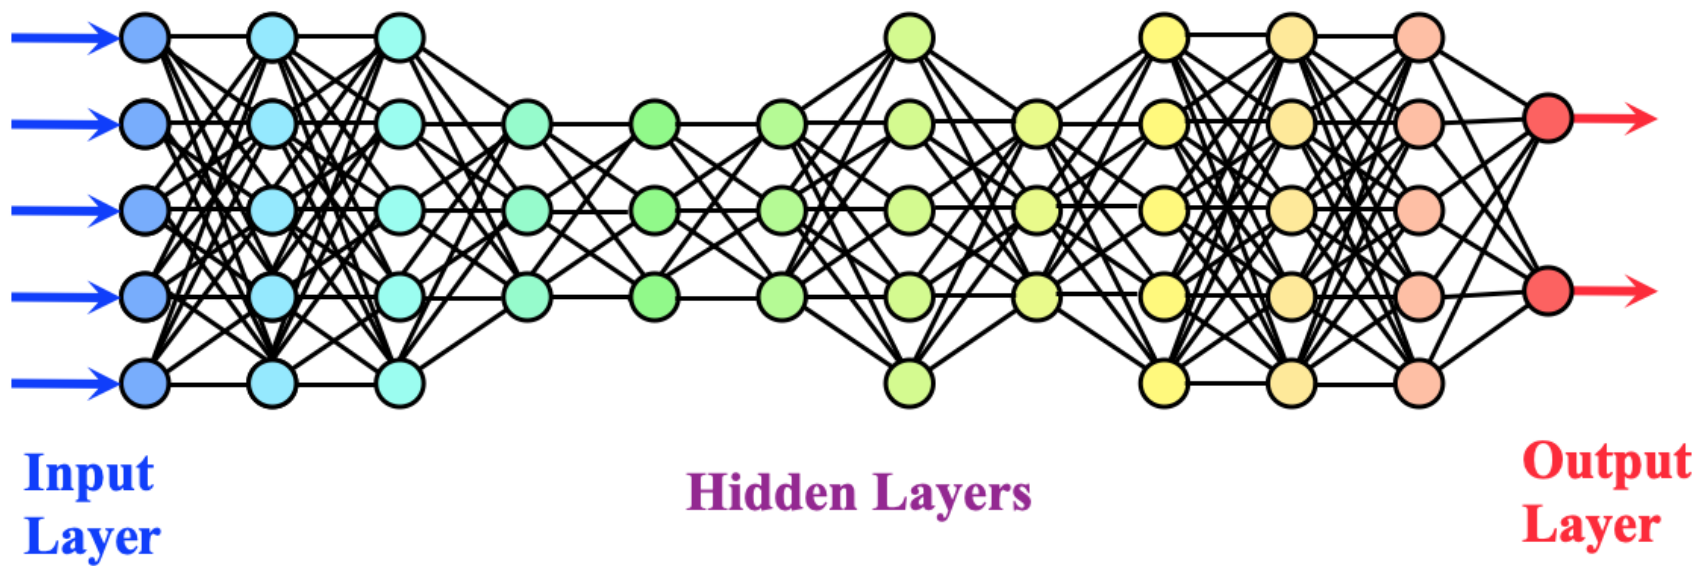
# Limitation of conventional CNNs

Very deep networks

Degradation problem

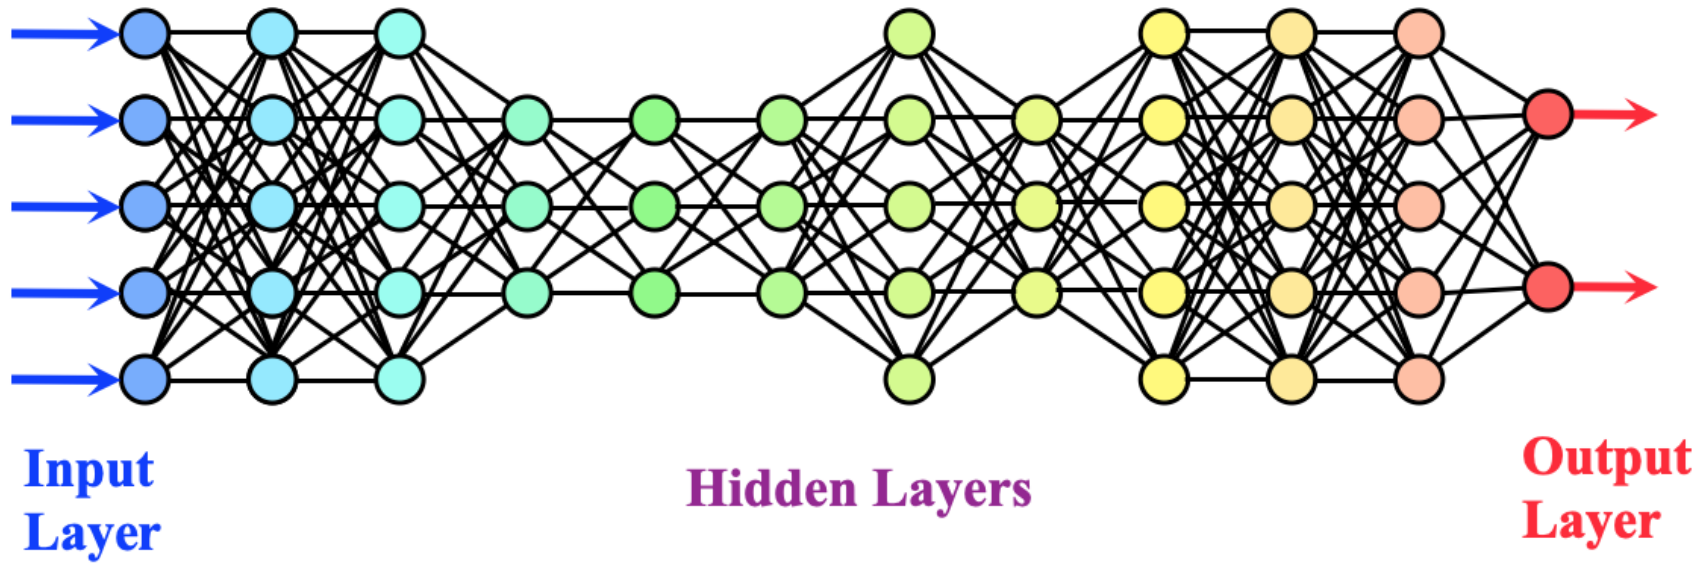


# Very Deep Networks



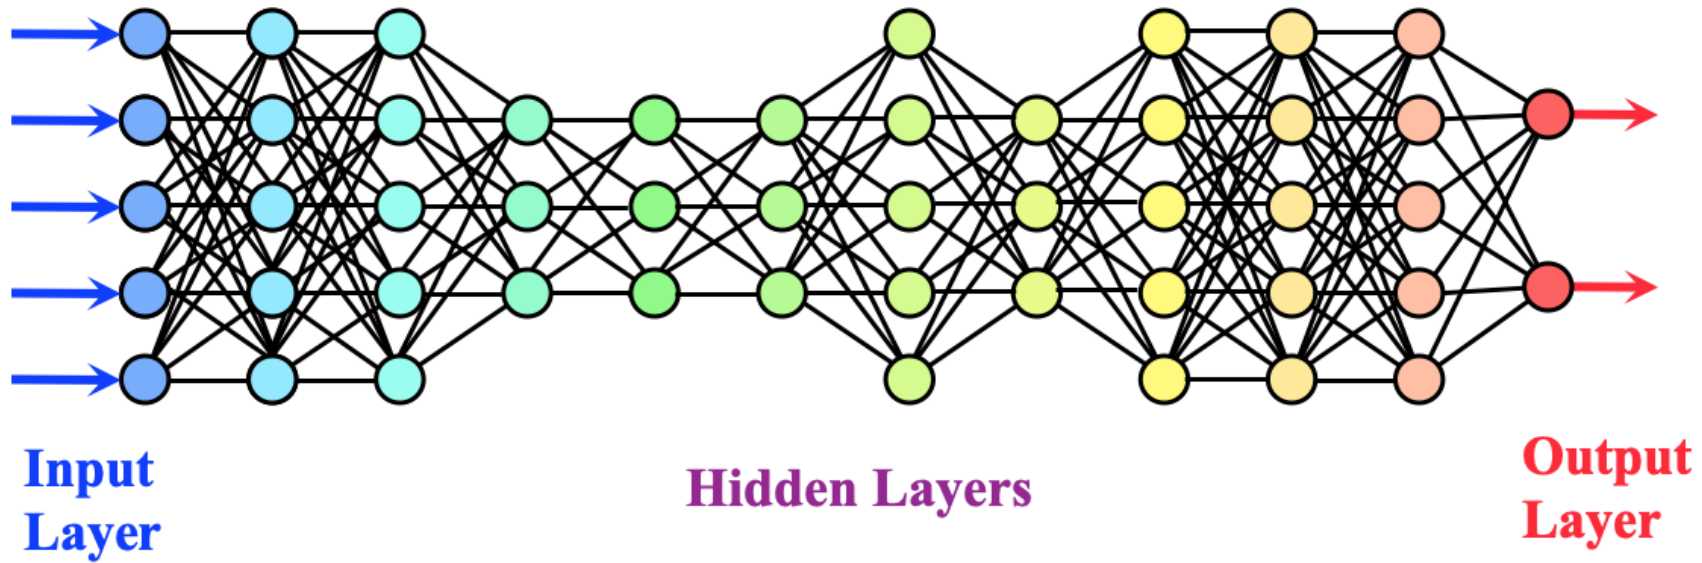


# Very Deep Networks





# Very Deep Networks



>100 layers

More layers → can process more complex patterns



# Degradation problem

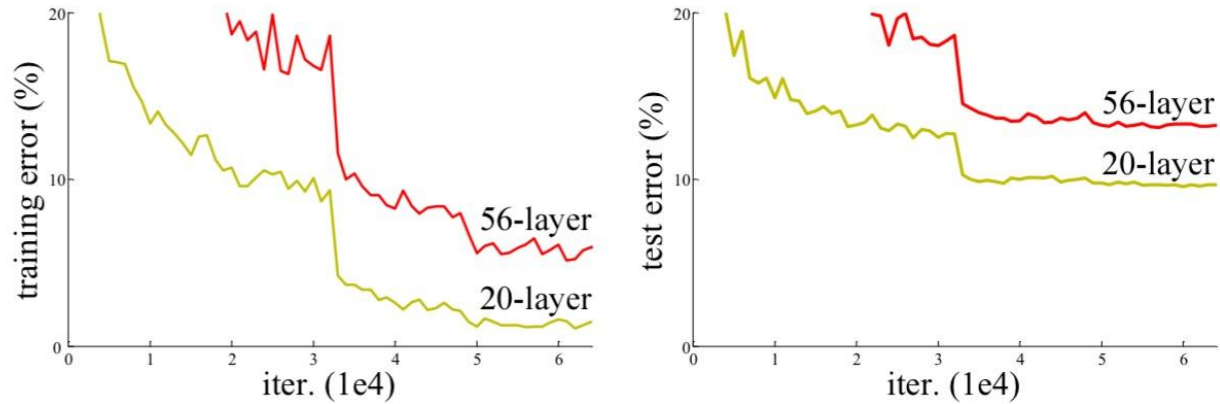


Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

A Review of the Evolution of Deep Learning Architectures and Comparison of their Performances for Histopathologic Cancer Detection



# Degradation problem

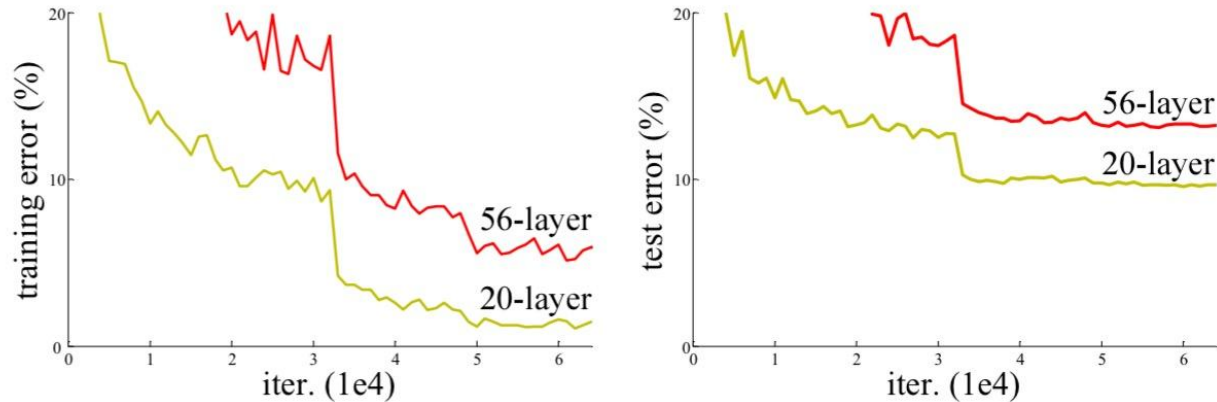


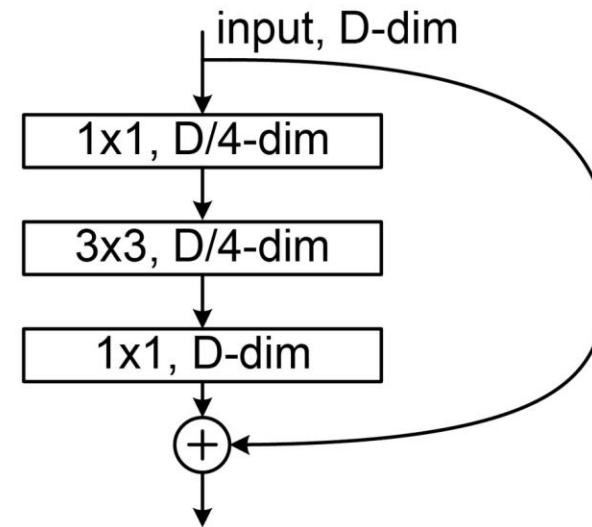
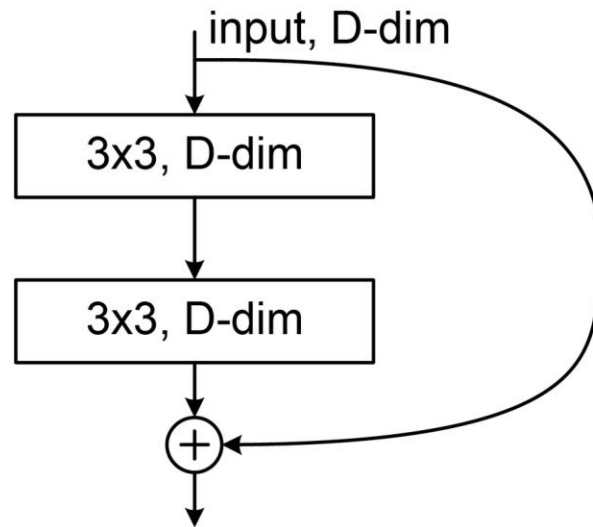
Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

Training / Testing accuracy start saturating faster than shallow networks



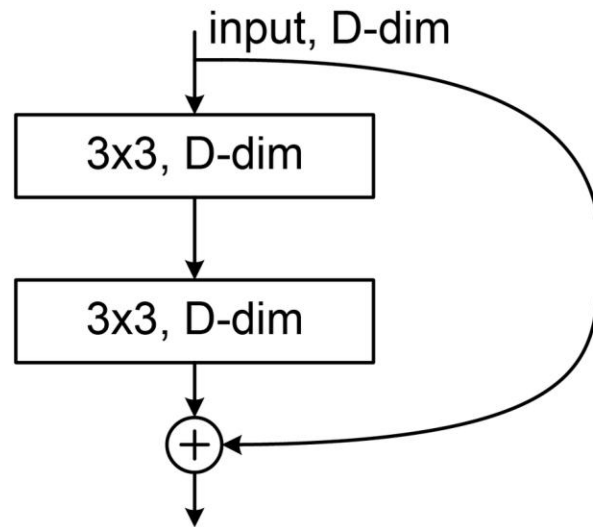


# How to train very deep networks

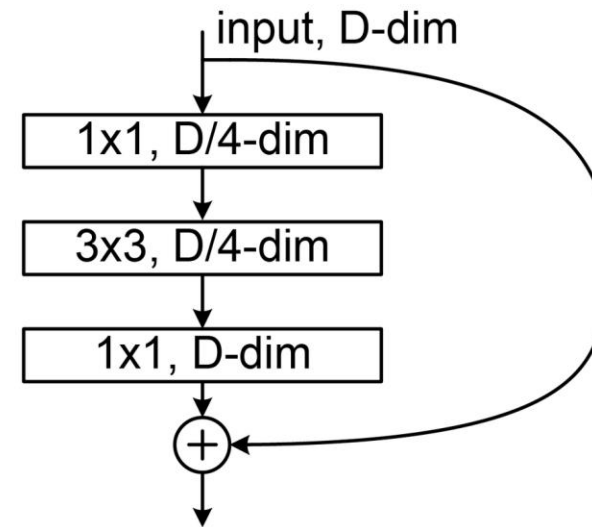




# How to train very deep networks



Skip connections  
(Residual block)



Data compression  
(Bottleneck layer)



# Additional CNN Blocks

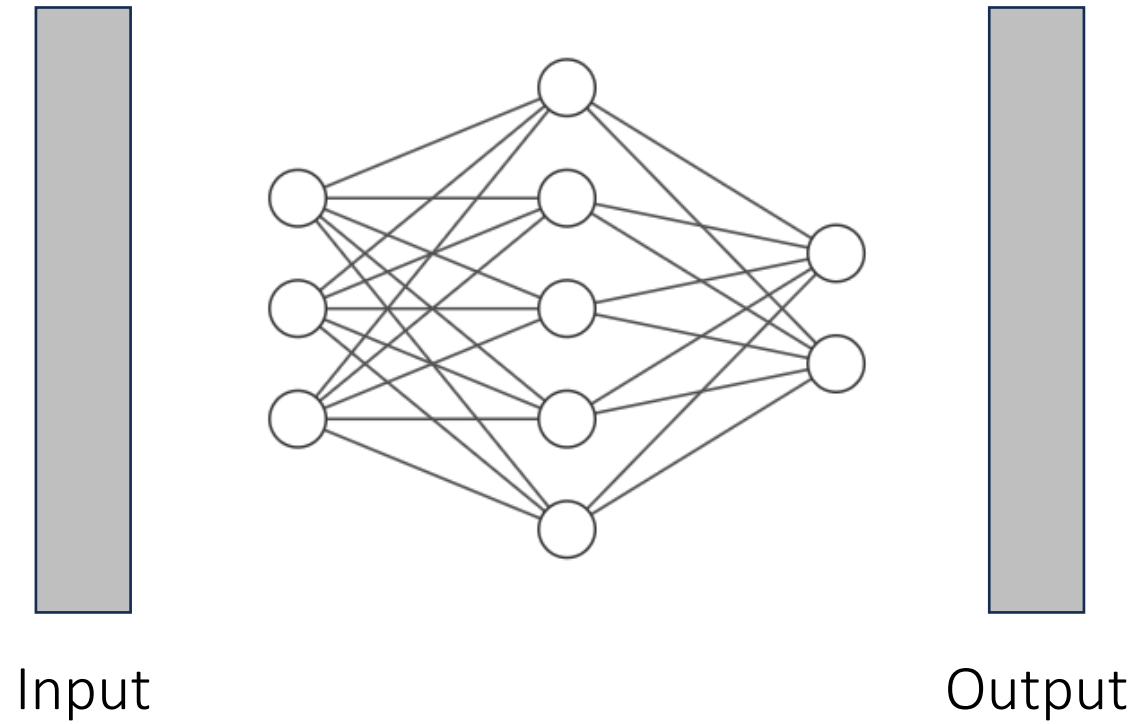
Residual blocks

Bottleneck blocks

Inception modules

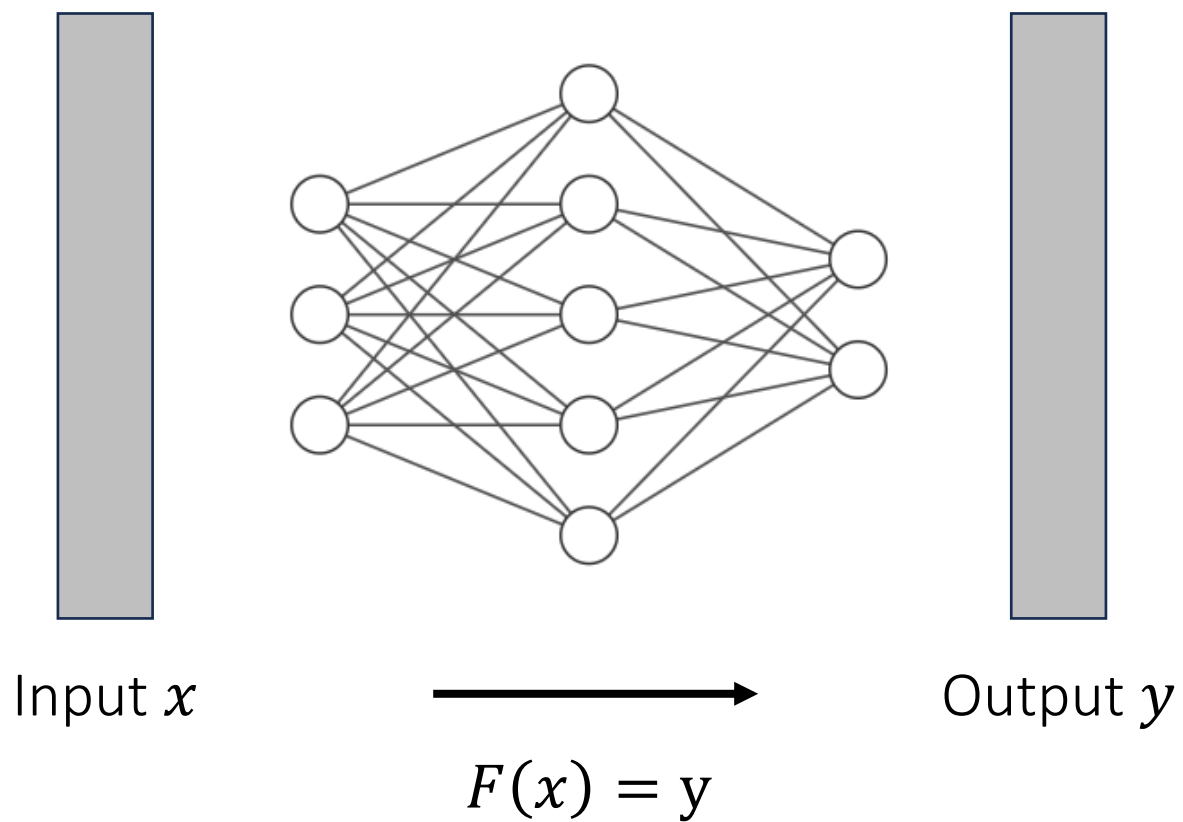


# Residual blocks



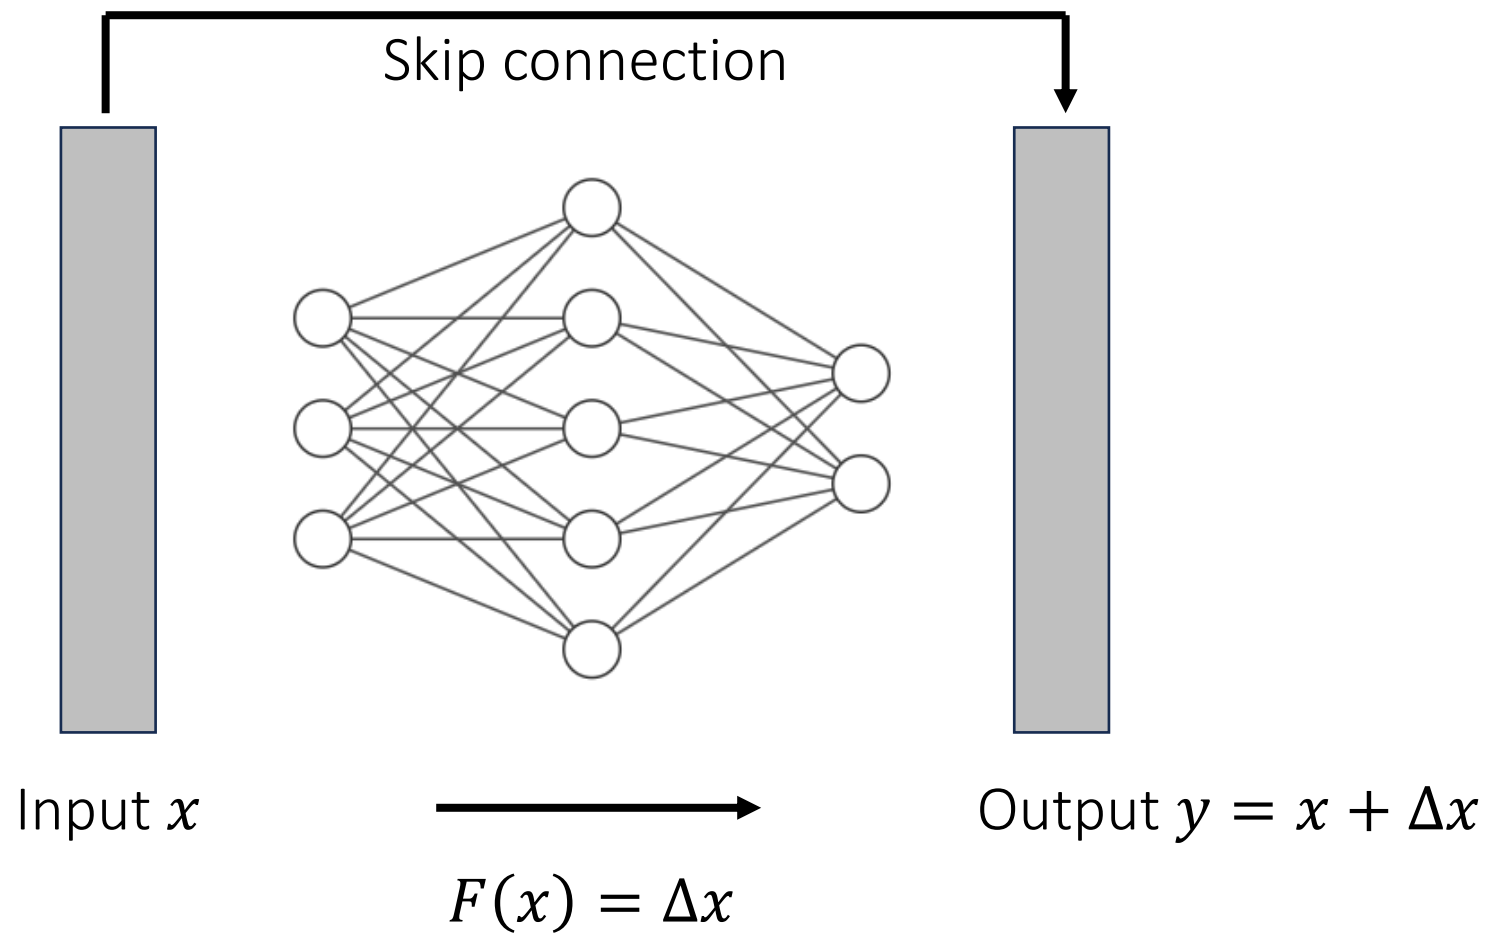


# Residual blocks



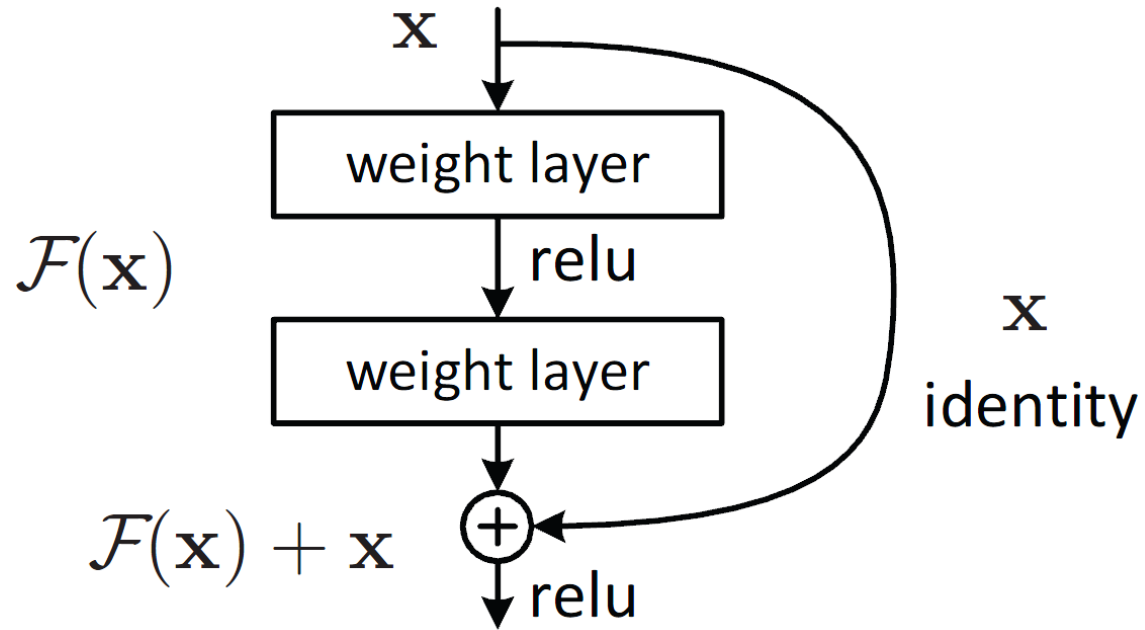


# Residual blocks





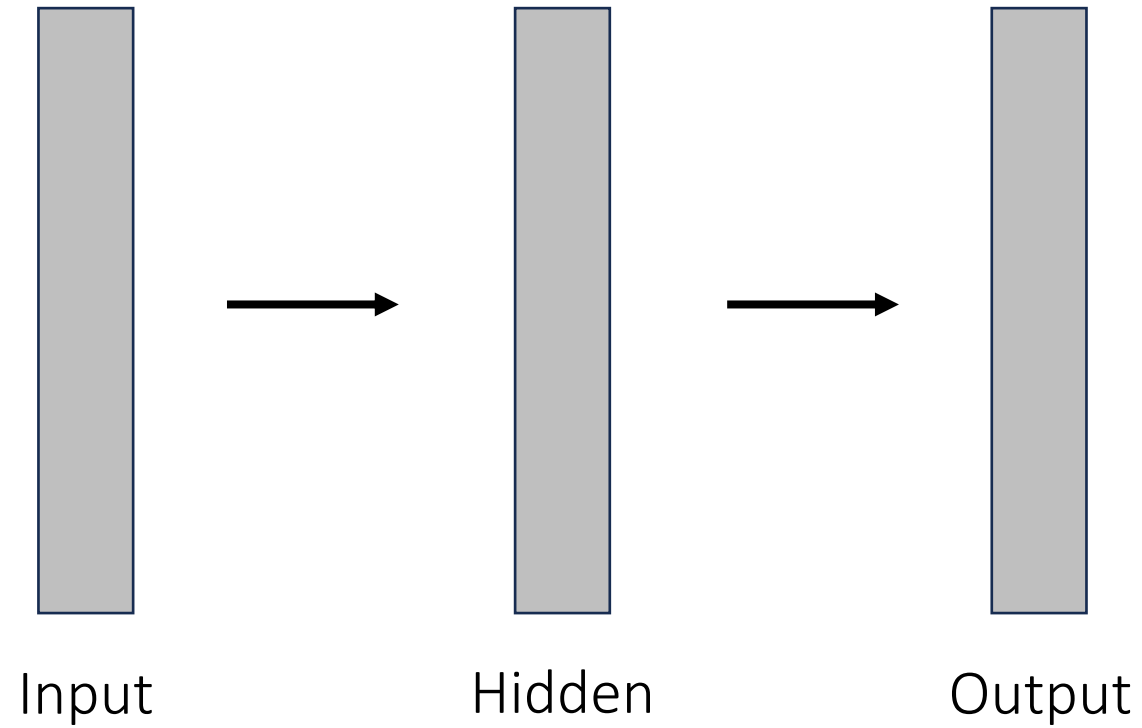
# Residual blocks



- Enable training deeper networks
- Prevents vanishing gradients
- Reduces overfitting



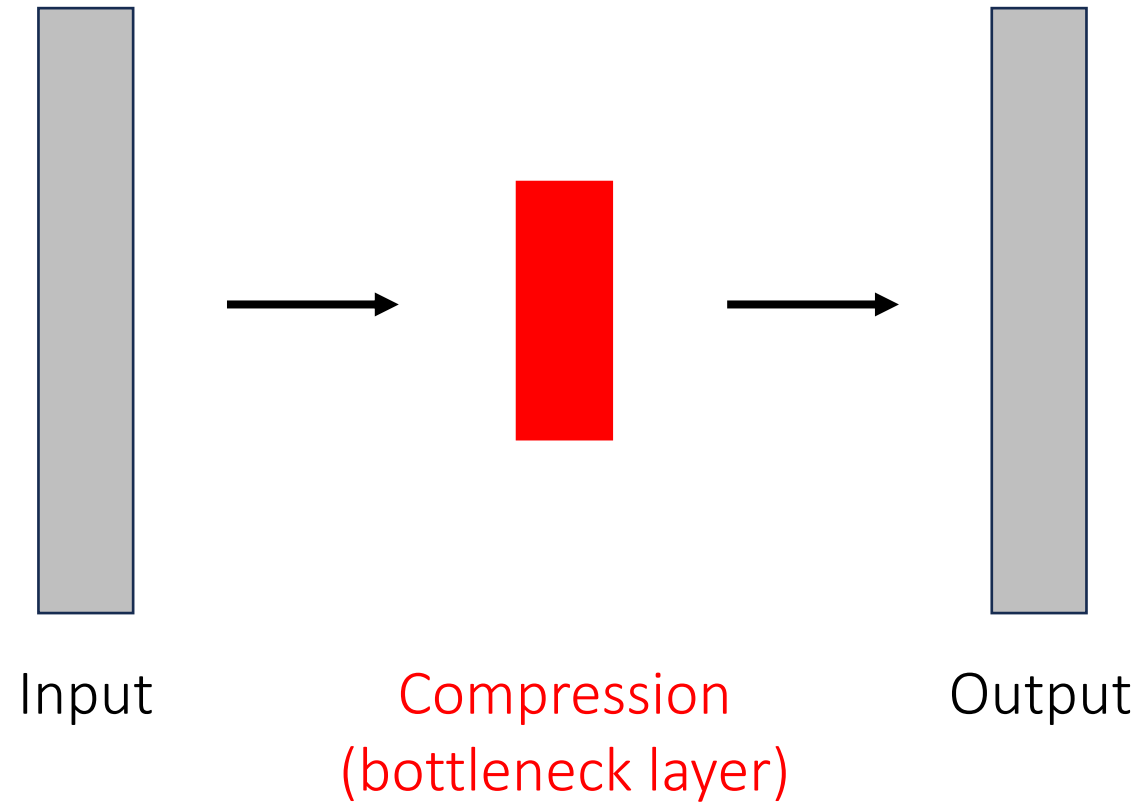
# Bottleneck layer





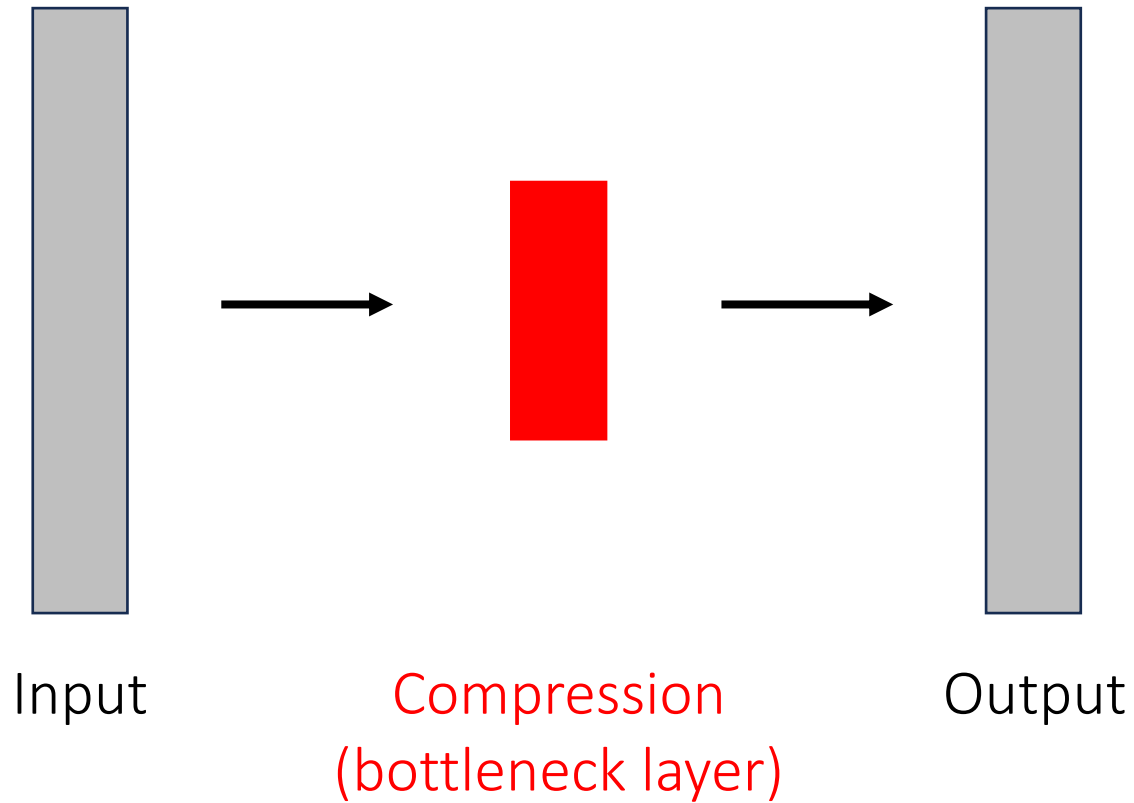


# Bottleneck layer





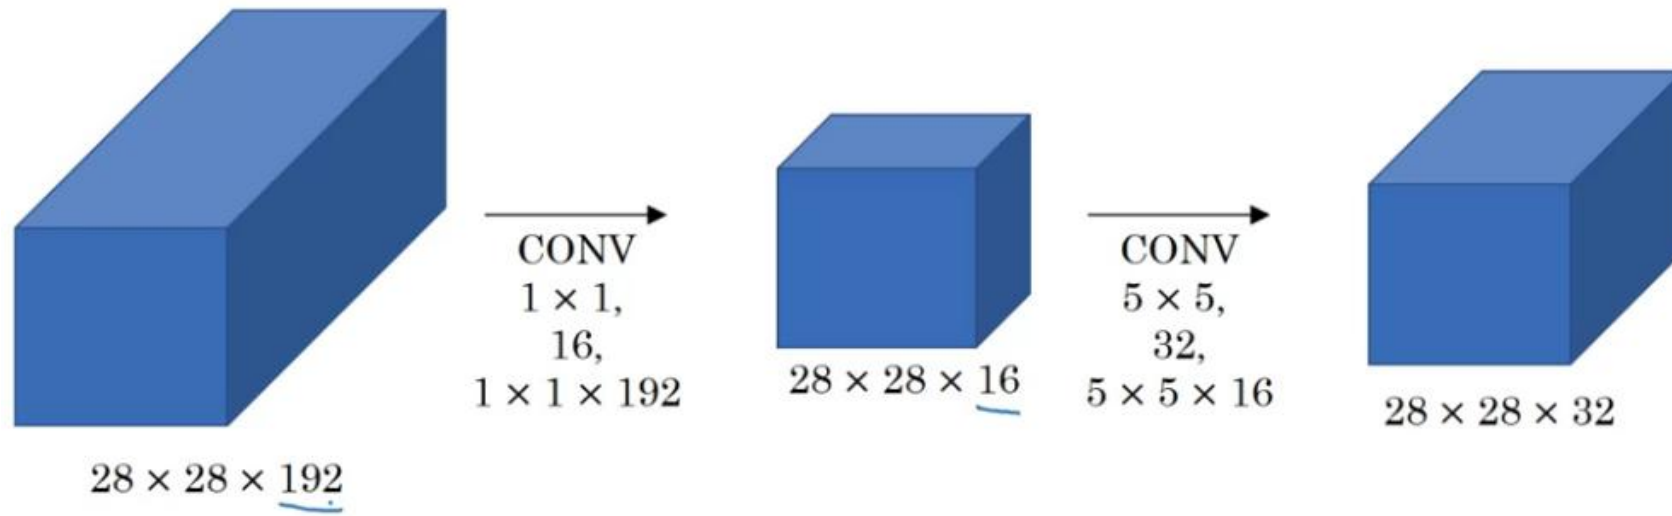
# Bottleneck layer



- Low dimensional representation
- More computationally efficient
- Reduces overfitting



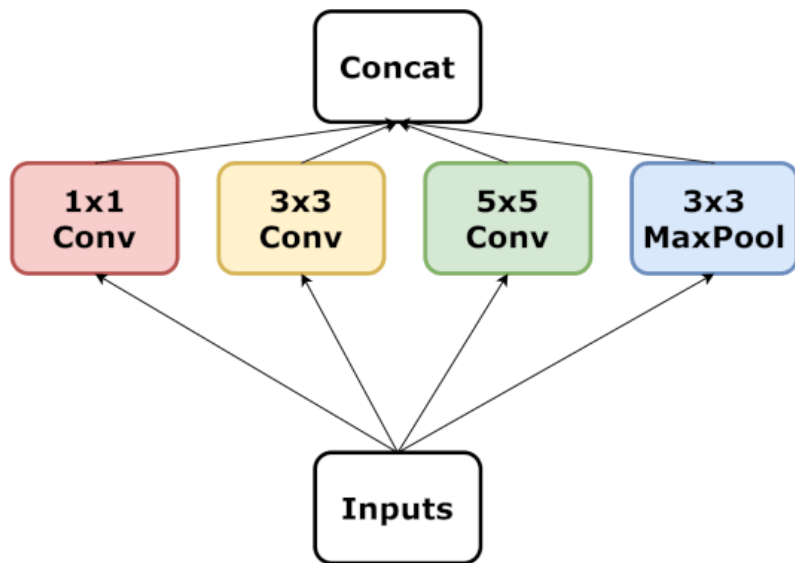
# Bottleneck layer in CNNs



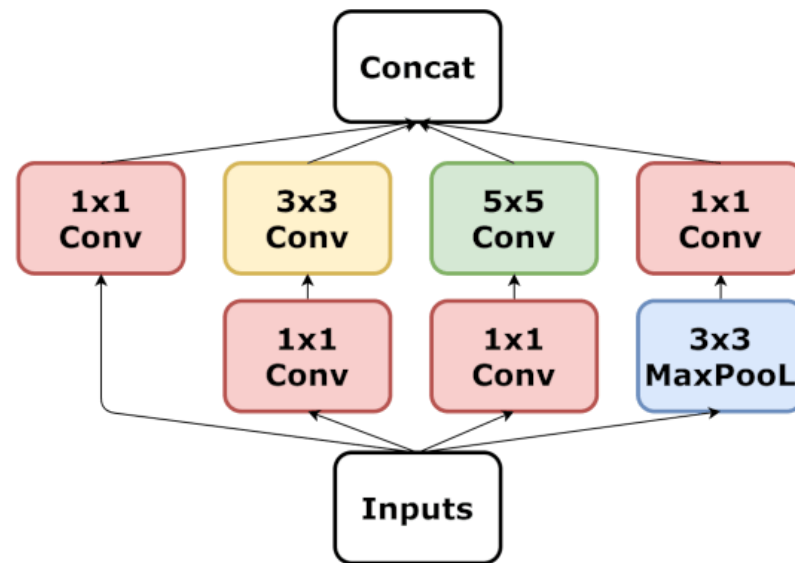
Compression  
(bottleneck layer)



# Inception module



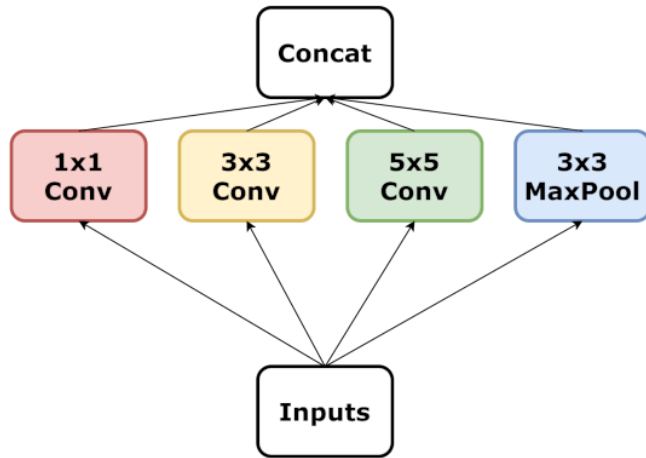
Inception Module



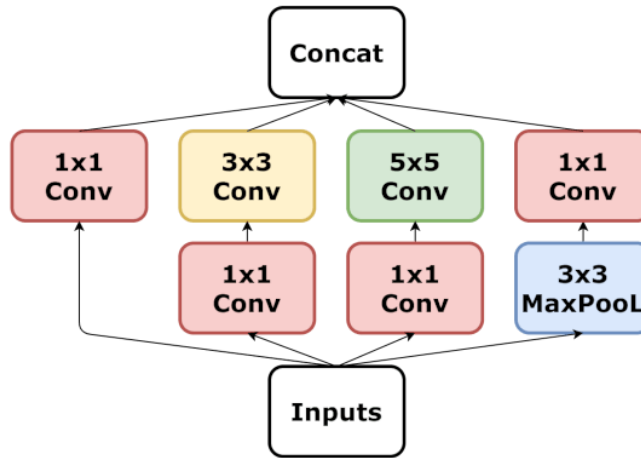
Inception Module with Dimension Reduction



# Inception module



Inception Module

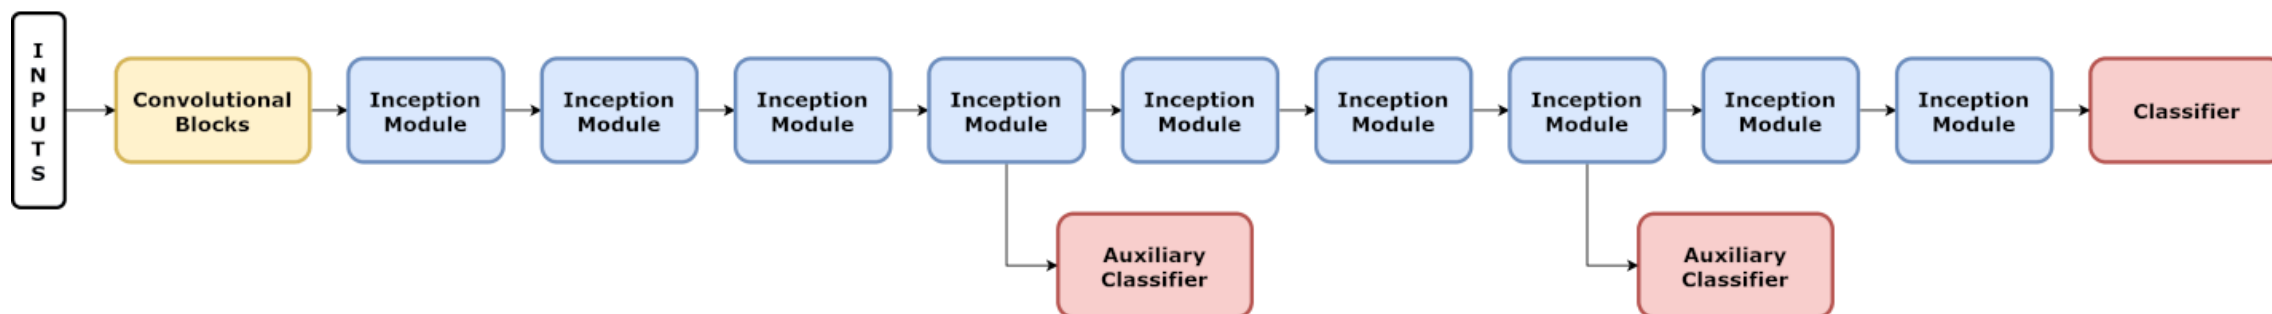


Inception Module with Dimension Reduction

- Parallel convolution operations
- Richer variety of features in a single output
- Can be used as standardized module



# Inception module



Inception V3 model



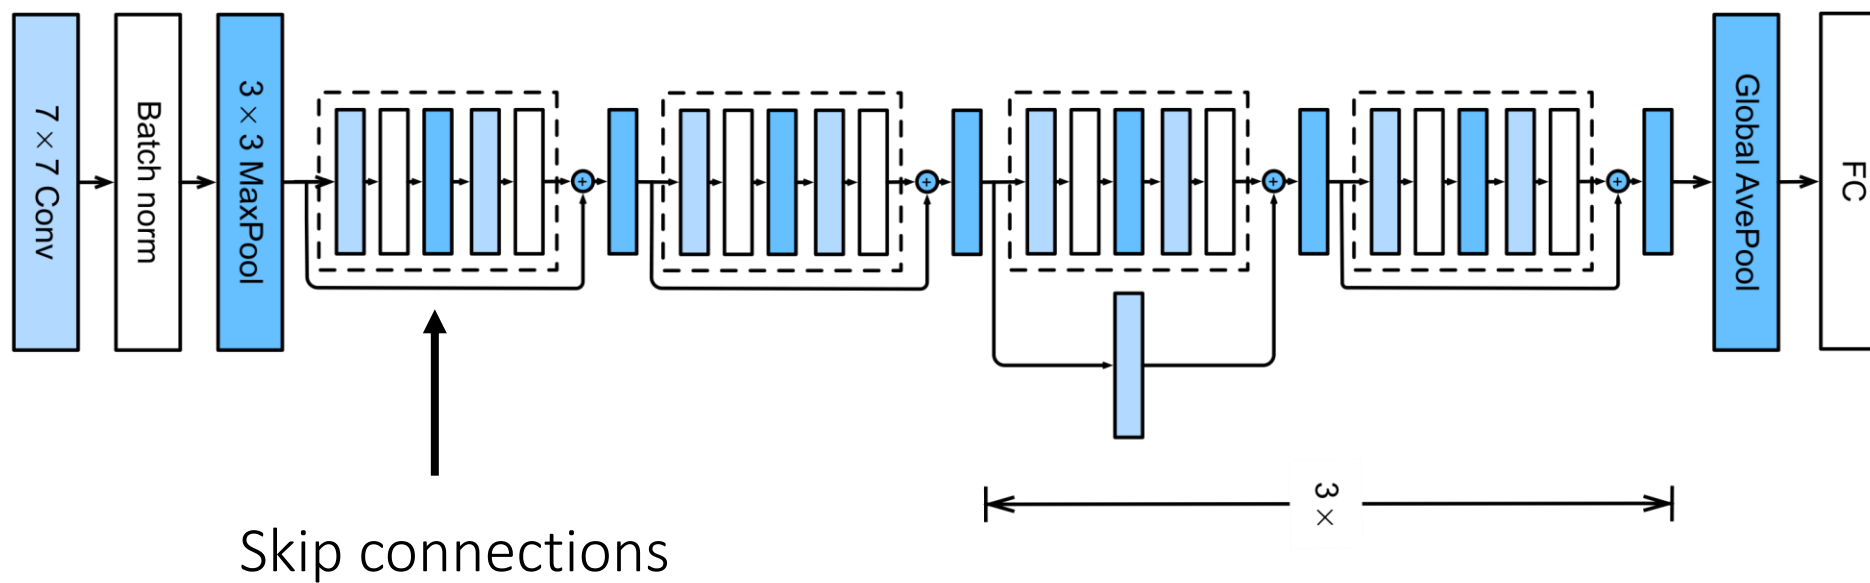
# Advanced CNN examples

ResNet-18

GoogLeNet



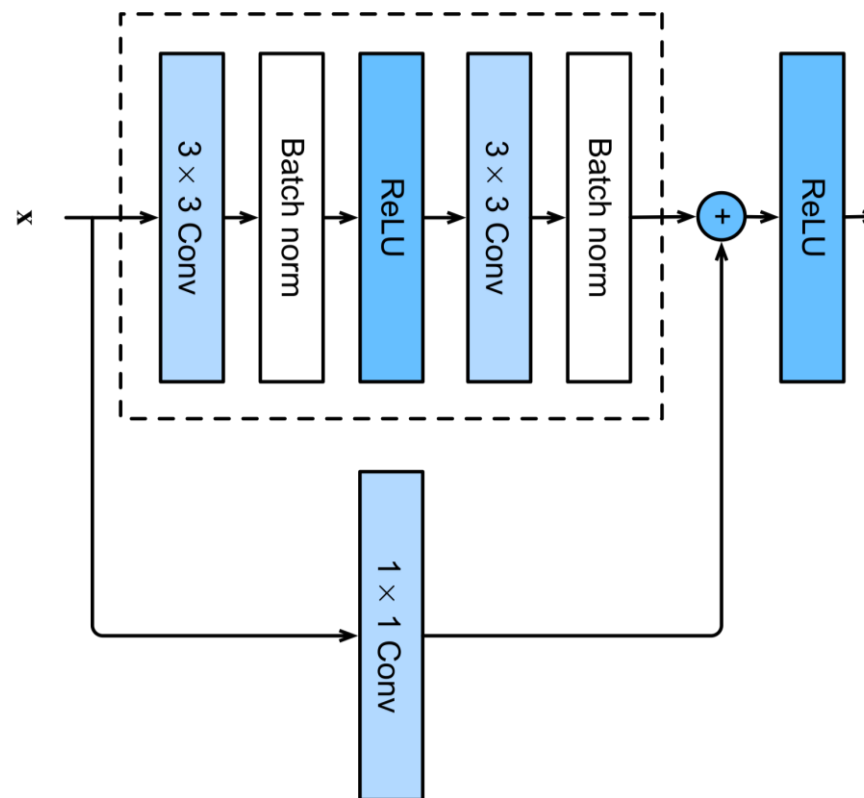
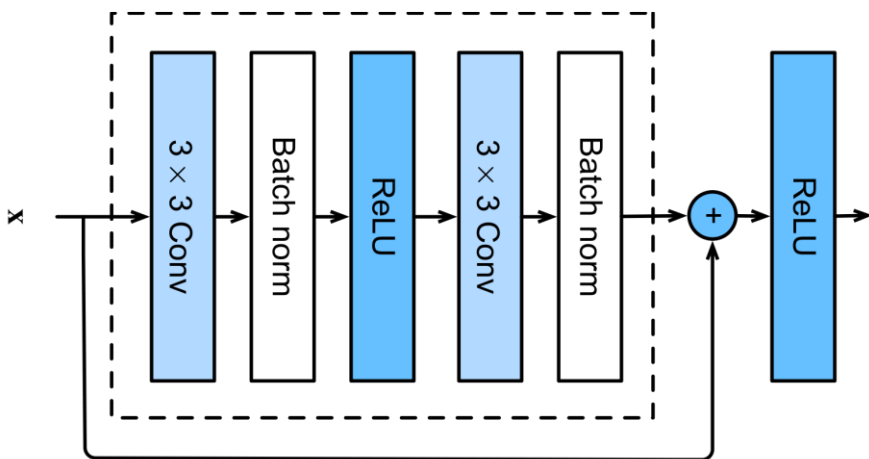
# ResNet-18







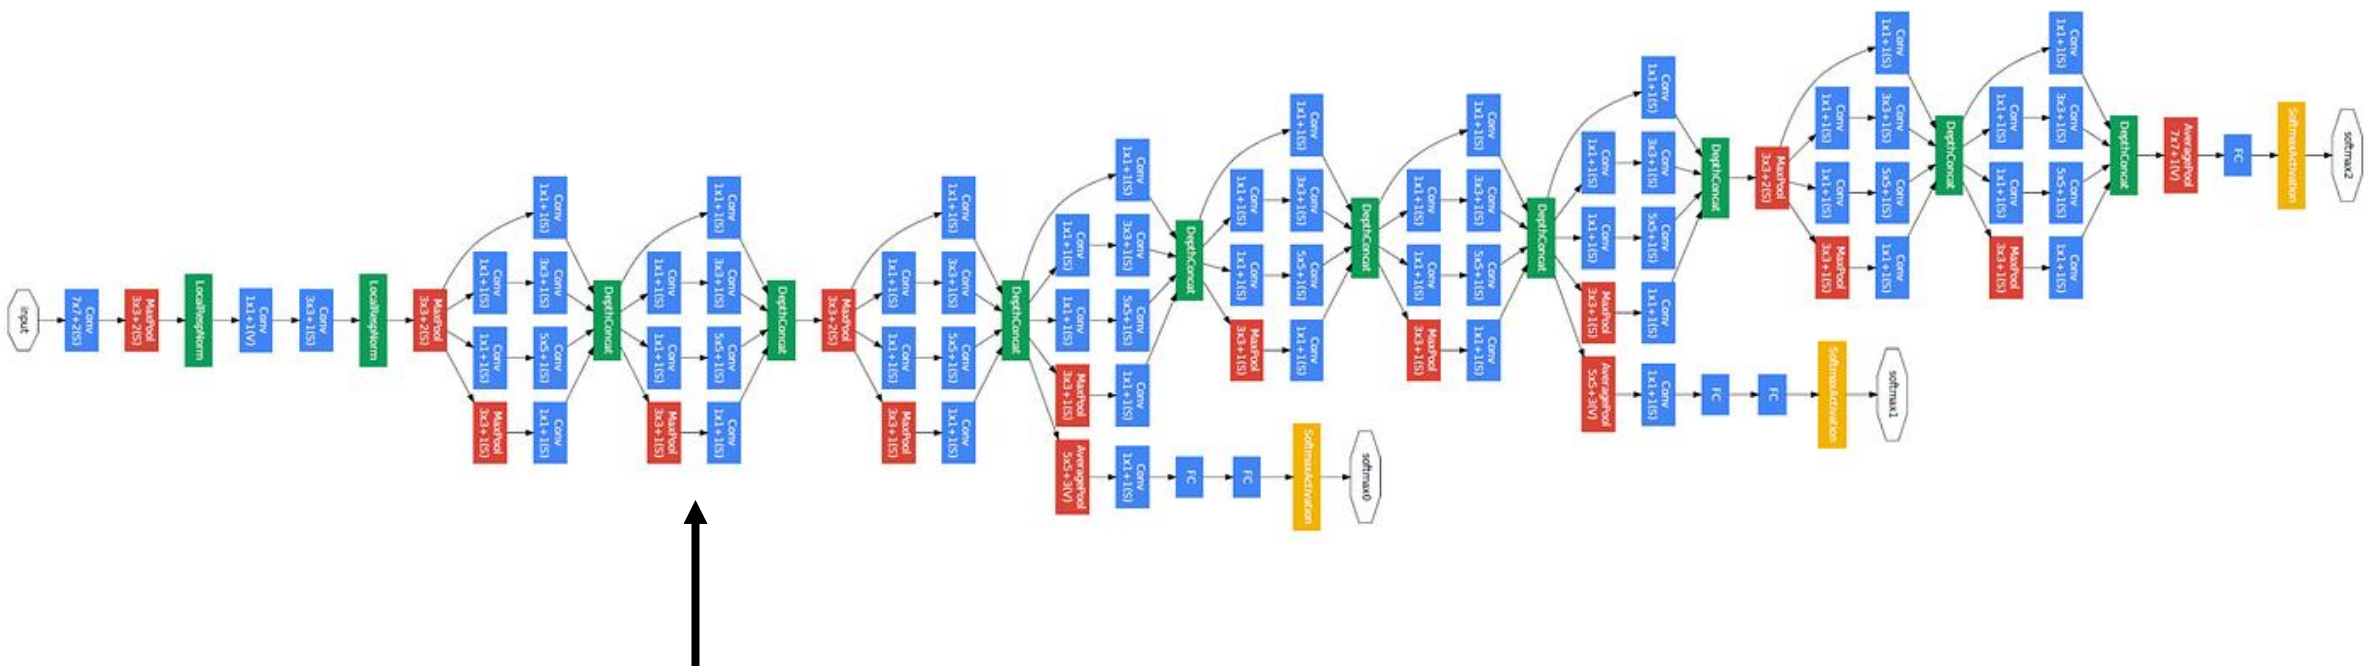
# ResNet-18



$1 \times 1$  conv for  
upsampling/downsampling



# GoogLeNet



Inception module



# ResNet Example

MNIST CLASSIFICATION



# Prepare Data

```
from torchvision import datasets, transforms

transform = transforms.Compose(
    [transforms.ToTensor(), transforms.Normalize((0.1307,), (0.3081,))]
)

train_dataset = datasets.MNIST(
    root="./data", train=True, download=True, transform=transform
)
test_dataset = datasets.MNIST(
    root="./data", train=False, download=True, transform=transform
)
```

Normalize the images

Load data from  
torchvision library



# Define Model

```
from torchvision.models import resnet18
```

```
class MNIST_ResNet_Classifier(torch.nn.Module):  
  
    def __init__(self, out_channels, kernel_size, stride, padding):  
  
        super(MNIST_ResNet_Classifier, self).__init__()  
  
        self.resnet = resnet18()  
  
        self.resnet.conv1 = torch.nn.Conv2d(1, out_channels = out_channels,  
                                             kernel_size=(kernel_size, kernel_size),  
                                             stride=(stride, stride),  
                                             padding=(padding, padding),  
                                             bias=False)  
  
        self.resnet.fc = torch.nn.Linear(self.resnet.fc.in_features, 10)  
  
    def forward(self, x):  
  
        return self.resnet(x)
```

Using pre-built model

Pre-built resnet-18 components  
can be edited to your liking



# Define Hyperparameters

```
MNIST_ResNet_Classifier = MNIST_ResNet_Classifier(out_channels = 64,  
                                                    kernel_size = 7,  
                                                    stride = 2,  
                                                    padding = 3)  
  
# Define Learning rate and epochs  
learning_rate = 0.001  
epochs = 10  
  
# Batch size for mini-batch gradient  
batchsize = 64  
  
# Using Adam as optimizer  
loss_func = torch.nn.CrossEntropyLoss()  
optimizer = torch.optim.Adam(MNIST_ResNet_Classifier.parameters(), lr=learning_rate)  
  
MNIST_ResNet_Classifier.cuda()
```

Initialize model

Initialize hyperparams

Using CE loss and  
Adam optimizer



# Identify Tracked Values

```
train_loss_list = []
```

Store training data loss



# Train Model

```
train_loader = torch.utils.data.DataLoader(  
    train_dataset, batch_size=batchsize, shuffle=True  
)  
  
test_loader = torch.utils.data.DataLoader(  
    test_dataset, batch_size=1000, shuffle=False  
)
```

Using dataloader to batch train and test dataset

```
for epoch in range(epochs):  
  
    for batch_idx, (train_input, train_target) in enumerate(train_loader):  
  
        optimizer.zero_grad()  
  
        pred = MNIST_ResNet_Classifier(train_input.cuda())  
  
        loss = loss_func(pred, train_target.cuda())  
  
        loss.backward()  
  
        optimizer.step()  
  
    print("Train Epoch: {} \tLoss: {:.6f}".format(epoch, loss.item()))
```

Training cycle

Print training loss each epoch





# Visualize and Evaluate Model

```
MNIST_ResNet_Classifier.eval()
test_loss = 0
correct = 0

with torch.no_grad():

    for test_input, test_target in test_loader:

        output = MNIST_ResNet_Classifier(test_input.cuda())

        test_loss += loss_func(output, test_target.cuda()).item()

        pred = output.argmax(dim=1, keepdim=True)

        correct += pred.eq(test_target.view_as(pred).cuda()).sum().item()

test_loss /= len(test_loader.dataset)

print("Test Accuracy: {}/{} ({:.0f}%)".format(correct, len(test_loader.dataset), 100.0 * correct / len(test_loader.dataset)))

Test Accuracy: 9921/10000 (99%)
```

99% accuracy within 10 epochs



# Additional UW Resources

UW Hayak supercomputer



# UW Hyak



- 30932 CPU cores
- 832 GPU (Turing, Ampere)
- Supports Python and Jupyter interface



# Access through UW RCC

## Hyak Access

### Applying for Access

1. Make sure you are eligible and have [\[joined RCC\]](#).
2. Review the [\[Hyak Prerequisites\]](#). Hyak is used through a linux command line. For many researchers, HPC resources are the first time they encounter this. This is fine! But you should spend a bit of time making sure you grasp the basics before moving to the supercomputer.
3. Read the [\[Hyak Documentation\]](#) in full. Further, you should consider attending an in-person [\[Hyak training session\]](#).
4. Read and agree to the RCC Hyak resource [\[terms of service\]](#).
5. Complete the [\[RCC Hyak skills assessment\]](#). Completion of this form serves as your application for access to Hyak and will trigger the review of all requirements. It may take up to 5 business days to grant access. If you have not heard back after 5 days, you may email RCC leadership or reach out to an officer on slack.

### Completing Set Up

Once you receive an email stating that your access has been granted, there's a few more steps:

1. Make sure two-factor authentication is set up. It almost certainly is, but you can check this at <https://itconnect.uw.edu/tools-services-support/access-authentication/2fa/>.
2. Be sure to [\[Subscribe to the Hyak mailing list\]](#) for reminders about monthly maintenance and training opportunity announcements.

- Join UW Research Computing Club
- Read Hyak Prereq / Documentation
- Take skill assessment for access
- Gain access to Hyak system



# Lab Assignment

CIFAR-10 Classification with ResNet-18



# CIFAR-10 Dataset

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



- 60000 32 x 32 colour images
- 10 classes
- 6000 per class
- 50000 train
- 10000 test



# CIFAR-10 Classification with ResNet



In this exercise, you will classify colour image (32 x 32 x 3) using your own variant of **ResNet-18**.

Prior to training your neural net, 1) Normalize the dataset and 2) Split the dataset into train/validation/test. You may also need to augment your dataset (e.g., flipping for better performance)

Feel free to tweak ResNet-18 architecture with your choices of **channels, kernel size, stride, padding** etc.

Your goal is to **achieve a testing accuracy of  $\geq 85\%$**  with no restrictions on epochs.

Demonstrate the performance of your model via plotting the **training loss, validation accuracy** and printing out the **testing accuracy**.

After your model has reached the goal, print the accuracy of each class similar to Lab 3. What is the class that your model performed best and worst?